

CSE-816(Final Project) - VQA MLOps

GitHub Repository: [GitHub Repo](#)
Youtube Project Demo: [YouTube Demo](#)

Team Members

- Trupti Khodwe(IMT2022007)
 - Rishabh Dixit (IMT2022051)
-

1. Project Description

This project implements a complete **MLOps (Machine Learning Operations) pipeline** for deploying a **Visual Question Answering (VQA) Agent**. The system automates the lifecycle of the ML model from code commit to production deployment using industry-standard DevOps tools.


Key capabilities include:

- **Hybrid Model Architecture:** Utilizes a lightweight text model (bert-tiny) for rapid CI simulation and a heavy VQA model (blip-vqa-capfilt-large) for production inference.
- **Efficient Artifact Management:** Implements "Model-as-Image" patterns and LoRA adapter integration to handle large weights efficiently.
- **Infrastructure-as-Code:** Uses Ansible for templated, dynamic Kubernetes deployments.
- **Observability & Metrics:** Full integration with the ELK Stack to track Latency and Model Confidence Scores.
- **Self-Healing Feedback Loop:** An automated watchdog service that triggers retraining pipelines if model quality drops below a threshold.
- **Scalability:** Horizontal Pod Autoscaling (HPA) to handle traffic bursts.

2. Architecture & Tech Stack

Component	Tool Used	Purpose
Source Control	GitHub	Version control for code and Infrastructure-as-Code manifests.
CI Automation	Jenkins	Automates testing, building, and orchestration.
Containerization	Docker	Builds optimized, layered images for the application and model artifacts.

Component	Tool Used	Purpose
Configuration	Ansible	Manages Kubernetes manifests via Jinja2 templating.
Orchestration	Minikube (K8s)	Local Kubernetes cluster for deployment.
Monitoring	ELK Stack	Visualizes logs and tracks quality_score & latency metrics.
Model	HuggingFace	Hosts the Base Model and Fine-Tuned LoRA Adapters.
Automation	Python Watchdog	Closes the MLOps loop by triggering Jenkins based on live metrics.



3. Prerequisites & Installation

A. Environment Setup (WSL/Linux)

Ensure the following tools are installed:

1. **Docker**.
2. **Minikube & Kubectl**: For local cluster management.
3. **Ansible**: pip install ansible kubernetes.
4. **Ngrok**: To expose Jenkins to GitHub Webhooks.

B. One-Time Data Setup

Since the dataset and models are massive (>3GB), we do not build them into the Docker image. We use **Host Volume Mounting**.

1. **Download Models locally**: Run the provided helper scripts to fetch weights into your data/ folder:
bash
python3 download_base.py # Saves to ./data/blip-large
python3 download_adapter.py # Saves to ./data/vqa-adapter
2. **Verify Structure**: Ensure ~/Final Project/data/ contains: blip-large/, vqa-adapter/, abo-images-small/, and VQA_Dataset.csv.

C. Cluster Initialization

Start Minikube with sufficient resources for the AI Model + ELK Stack.

```
minikube start
```

D. Permission Fixes (Critical for Jenkins)

Allow the Jenkins user to access the Minikube configuration files:

```
sudo chmod +x /home/$USER  
sudo chmod -R 777 /home/$USER/.minikube  
sudo chmod -R 777 /home/$USER/.kube  
sudo chmod a+x /home/user
```

4. Critical Manual Commands (The "Glue")

These commands must be running in separate terminal windows for the project to function.

1. Data Mounting (The Bridge)

Connects your local hard drive (Dataset + Models) to the Minikube VM.

```
# Run in Project Root  
minikube mount ./data:/data
```

- Why: Allows the Pods to read the 3GB dataset without downloading it.

2. Application Tunnel (Accessing the AI)

Exposes the VQA Agent to localhost:5000.

```
kubectl port-forward svc/vqa-service 5000:80 --address 0.0.0.0
```

- Why: Minikube networks are isolated. This creates a direct pipe for curl.

3. Kibana Tunnel (Accessing Logs)

Exposes the Monitoring Dashboard to localhost:5601.

```
kubectl port-forward -n logging svc/kibana 5601:5601 --address 0.0.0.0
```

4. Elasticsearch Tunnel (For Watchdog Script)

Exposes the ES API for the feedback loop script.

```
kubectl port-forward -n logging svc/elasticsearch 9200:9200 --address 0.0.0.0
```

5. The CI/CD Pipeline Workflow

The Jenkinsfile defines a 6-stage pipeline triggered automatically by a Git Push.

- Checkout: Pulls the latest code from GitHub.
- CI: LoRA Simulation: Runs train_mini.py to simulate the training process and generate dummy artifacts.
- Build Docker Images: Builds:
- vqa-app: The Flask Application.
- vqa-model: A layered container for model weights.
- CI: Model Evaluation: Runs evaluate_model.py inside a container. It loads the Real Model from the mounted disk and tests inference on a subset of real data.
- CD: Push to Registry: Pushes tagged images to Docker Hub (rishabh720/vqa-app).
- CD: Deploy to Kubernetes:
- Loads images directly into Minikube (bypassing internet/bandwidth issues).
- Uses Ansible Roles (playbooks/roles/vqa_deploy) to render deployment.yaml.j2 with the new tag.
- Applies the deployment + HPA + ELK Stack.



6. Advanced Features & Innovation

A. Modular Ansible Roles

- Instead of a monolithic playbook, we use a modular design:
- Role: playbooks/roles/vqa_deploy
- Function: Separates template rendering and manifest application logic, ensuring scalability and maintainability.

B. Quality & Performance Metrics

- The application calculates and logs specific metrics for every inference request:
- quality_score: The model's confidence probability (Softmax) in its generated answer.
- latency_seconds: The exact time taken to process the image and generate text.

C. Automated Feedback Loop (Self-Healing)

- A watchdog script (monitor_feedback_loop.py) runs alongside the cluster.
- It queries Elasticsearch for recent quality_score metrics.
- If the average score drops below a threshold (e.g., 0.75).
- It automatically triggers the Jenkins Pipeline via API to retrain/redeploy the model.

The screenshot displays three panels illustrating the system architecture and configuration:

- Top Panel:** A terminal window showing the command `kubectl get hpa` and its output. The output shows a single HPA named "vqa-agent-hpa" which is targeting the "Deployment/vqa-agent" and has a minimum of 1 pod, a maximum of 5 pods, and a current replica count of 1. It was created 46 hours ago.
- Middle Left Panel:** A terminal window showing the content of the `deploy.yml` Ansible playbook. The playbook defines a single host at `localhost` using `local` connection, gathers facts, and deploys a role named `vqa_deploy`.
- Middle Right Panel:** A file browser interface showing the directory structure of the Ansible role `roles/vqa_deploy`. It contains `tasks/main.yml`, `templates/deployment.yaml.j2`, and `playbooks/deploy.yml`.

```
○ rishabh@Rishabh:~/Final Project$ python3 monitor_feedback_loop.py
--- Watchdog Started ---

[23:08:18] ● Watchdog: Checking Model Health...
    -> Latest Score: 0.6514
    ▲ALERT: Score 0.6514 < 0.75
    🚀 TRIGGERING RETRAINING PIPELINE...
    -> CSRF Crumb obtained.
    ✅ Pipeline Triggered Successfully!
    ❌ Cooling down for 30 seconds...[
```

7. How to Test & Demo

A. Trigger the Pipeline

- Make a change to the code.
- Push to GitHub: git push origin main.
- Watch the build succeed in Jenkins.

B. Verify "Real Intelligence"

With the Port Forward (Step 4.2) active, run this request:

```
curl -X POST http://localhost:5000/predict \
-H "Content-Type: application/json" \
-d'{
  "image_path": "14/14db355d.jpg",
  "question": "what is this item?"
}'
```

Expected Output:

```
{
  "answer": "sneakers",
  "event": "inference",
  "latency_seconds": 90.5573,
  "quality_score": 0.6514,
  "question": "what is this item?"}
```


```

● rishabh@Rishabh:~/Final Project$ curl -X POST http://localhost:5000/predict -H "Content-Type: application/json"
"      -d '{
        "image_path": "14/14db355d.jpg",
        "question": "what is this item?"
    }'
{"answer":"sneakers","event":"inference","latency_seconds":91.529,"quality_score":0.6514,"question":"what is this item?"}
○ rishabh@Rishabh:~/Final Project$ █

```

C. Verify Monitoring

- Open Browser: <http://localhost:5601>.
- Go to Discover.
- Search for: "LOG_EVENT".
- Evidence: You will see the structured JSON logs containing the quality_score and latency.



D. Verify Feedback Loop

Run the watchdog script manually to simulate the automated check:
`python3 monitor_feedback_loop.py`

- Output: It will detect the score (0.65) is below threshold (0.75) and print: 🚀 TRIGGERING RETRAINING PIPELINE....
- Result: A new build will appear in Jenkins.

```
○ rishabh@Rishabh:~/Final Project$ python3 monitor_feedback_loop.py
--- Watchdog Started ---

[23:08:18] 🌐 Watchdog: Checking Model Health...
-> Latest Score: 0.6514
⚠️ ALERT: Score 0.6514 < 0.75
⚡ TRIGGERING RETRAINING PIPELINE...
-> CSRF Crumb obtained.
✅ Pipeline Triggered Successfully!
⚡ Cooling down for 30 seconds...!
```

Builds

Filter

Pending

...

#60

Finished waiting

Today

#59 10:53 PM